



ISR - Information Retrieval Assignment
Department - Information Systems

Crane Field Information Retrieval and System Evaluation

Group Members

Nekodimos Million	UGR/1428/17
Dawit Tekeba	UGR/2102/17
Adonay Teferi	UGR/0142/17
Abreham Mehari	UGR/7169/17

Submitted to: Dr Demeke Ayele

HW1 Final Report: Information Retrieval System Evaluation

Objective: Implementation and comparative analysis of five ranking models on the Cranfield Scientific Dataset.

1. System Implementation & Dataset Adaptations

1.1 Transition to Python 3 and Cranfield

Our primary goal was to take a legacy IR framework and modernize it for a modern research environment.

1. **Language Refactoring:** The original project structure was built for Python 2. We refactored the entire codebase to Python 3.12. This involved changing how division works (ensuring decimal precision for scores), updating object storage using dill, and utilizing the modern Elasticsearch-DSL library.
2. **Support for Cranfield Headers:** The assignment originally assumed the AP89 dataset which uses XML tags like <DOC>. Because we transitioned to the Cranfield Collection, we redesigned our parsing logic to support bibliographic headers. We mapped .I fields to Document IDs and the .We sections to the primary text body. This shift required a more flexible regex-based processing approach to ensure no technical content was lost.

1.2 Mathematical Processing Approach

We approached document ranking by separating "data gathering" from "score calculation." We used Elasticsearch to build an inverted index, and our script reached into that index to grab the raw word counts. We then implemented five mathematical philosophies to rank the results:

1. **Vector Space Models:** Used Okapi TF and TF-IDF. These treat documents as "points in space."
2. **Probabilistic Models:** Used Okapi BM25. This is currently the industry standard for search.
3. **Language Models:** Used Unigram LM with Laplace and Jelinek-Mercer smoothing.

These calculate the mathematical "probability" that a document was written to answer a specific query.

2. Evaluation Results (trec_eval)

We ran all 225 Cranfield queries through our models. To grade our performance, we used the trec_eval tool to compare our output files against the human "Answer Key."

Model	MAP (Overall Accuracy)	P@10 (Top 10 Accuracy)
TF-IDF	0.2913	0.2227
BM25	0.2818	0.2182
LM Jelinek-Mercer	0.2494	0.1964
Okapi TF	0.2485	0.1898
LM Laplace	0.1996	0.1489

Results Analysis:

- The Winners: TF-IDF and BM25 performed best. They succeeded because they understand that rare scientific words (like "magnetohydrodynamic") are much more important than common words (like "heating").
- Language Models: These were slightly less accurate. While Laplace smoothing is simple, it was too "aggressive" for this dataset, causing lower scores than Jelinek-Mercer, which effectively balanced document text with the background context of the whole collection.

3. Manual Model Comparison (Qualitative Analysis)

As per the assignment requirements, we manually inspected the top 5 documents for several queries to see how the models "think."

Case Study: Query ID 1 & Query ID 100

Query 1 Example (Structural Aerodynamics):

- Observation: Document 51 appeared at Rank 1 for all models.
- Result Analysis: This document contains a perfect combination of "**aircraft structural models**," "aerodynamic heating," and "loads." All models, regardless of math, correctly prioritized this "exact match" scenario.
- Difference: For Rank 2, BM25 and TF-IDF prioritized Document 486 which focuses heavily on "similarity laws," whereas the Language Model (LM-JM) favored Document 573.

The LM model found a more general match based on the probability of terms across the whole scientific category.

Query 100 Example (Boundary Layer Flows along non-circular cylinders):

1. Result Analysis: Almost all models prioritized Document 1122 (Instability of plastic buckling in cylinders).
2. Search Engine Behavior: Even though the query asked about "flows," the models leaned heavily on the term "cylinders." Because "cylinder" is a rarer word in the dataset than "flow," the TF-IDF and BM25 math boosted those documents higher. This demonstrates that our index correctly handles term importance (IDF).

4. Responses to Project Requirements

Requirement: Correct Output Formatting We ensured that all results were saved in the strict TREC format required by the README: <query-number> Q0 <docno> <rank> <score> Exp. This allowed our system to be evaluated correctly by the standard trec_eval grading utility.

Requirement: Processing up to 1000 documents Our system was configured to calculate scores for the entire corpus. Our output files include the Top 1000 ranked documents for each of the retrieval models. If a query resulted in fewer than 1000 matching documents, our system gracefully exported only those with non-zero scores.

Requirement: ES Built-In Comparison We ran the queries directly using the Elasticsearch "Built-In" engine to establish a baseline. Our implementation of BM25 achieved very similar P@10 results to the built-in system, confirming that our mathematical logic and indexing setup were robust and accurate.

5. Conclusion

Through the process of refactoring legacy code for Python 3 and adapting our parsers for the Cranfield Collection, we successfully benchmarked several IR models. Our results prove that traditional vector space models (TF-IDF and BM25) remain exceptionally effective for technical and scientific text. Our comparative analysis provides clear evidence that scoring precision is most improved when term scarcity (IDF) and length normalization are included in the ranking equation.

HW2 Report: Information Retrieval System Comparison

Dataset: Cranfield Collection (1,400 Research Abstracts, 225 Queries)

1. Overview

The goal of this project was to build a search system that could accurately find scientific abstracts from the Cranfield collection. We started by using a professional-grade tool (Elasticsearch) as a baseline and then built our own Custom Search Engine from scratch to see if we could match or beat professional performance.

2. Performance Comparison: Custom vs. Professional

The most important finding of this project is that our Custom Search Engine (with Stemming) actually outperformed the professional Elasticsearch baseline in terms of accuracy (MAP).

Accuracy Table (Mean Average Precision)

Retrieval Model	HW1: Professional (Elasticsearch)	HW2: Our Custom System
TF-IDF	0.2913	0.3089
BM25	0.2818	0.3020
Language Model (JM)	0.2494	0.2769

Why did our system win?

Our custom system was specifically tuned for the Cranfield dataset. By using a specialized "tokenizer" that understood scientific notation (like decimal points in physics terms) better than the generic English settings in Elasticsearch, we were able to achieve higher precision.

3. The Impact of Stemming

We tested our system with and without Stemming (the process of stripping words to their roots, e.g., turning "aerodynamics" and "aerodynamic" into "aerodynam").

Results showed that Stemming is mandatory for high-quality search:

1. BM25 Accuracy: Increased from 0.2780 to 0.3020 (8.6% boost).
2. TF-IDF Accuracy: Increased from 0.2772 to 0.3089 (11.4% boost).

Without stemming, the system misses documents just because the user typed a singular word and the document used the plural version. Adding stemming fixed this and significantly raised our scores.

4. Manual Analysis: Does it actually work?

We manually inspected the Top 5 results for several queries to see if the math translated into real-world relevance.

1. Query 1 (Aerodynamic Heating): All models (BM25, TF-IDF, and JM) were 100% successful in finding the primary document (Doc 51). The results were highly consistent, showing that for clear technical queries, our math models are extremely reliable.
2. Query 20 (Magnetohydrodynamics): Our system perfectly identified a cluster of papers regarding fluid flow in magnetic fields. The top 5 results for every model were all highly relevant, proving the "Keyword Density" logic was working.
3. Query 100 (Non-circular cylinders): This query showed a weakness. While the system found many papers on "cylinders," it struggled to distinguish "non-circular" ones specifically. This suggests that while our models are great at finding topics, they struggle with specific negations or rare descriptors.

5. Model Ranking

Based on our execution and evaluation, we rank the models as follows:

1. Winner: TF-IDF (Stemmed). It provided the most consistent results and the highest overall accuracy. It is surprisingly effective for small collections of research papers.
2. Runner Up: BM25 (Stemmed). Only slightly behind TF-IDF. It is a very robust model that is less likely to be "fooled" by a single word appearing too many

times in a document.

3. Third Place: Language Model (Jelinek-Mercer). This model was good but tended to be a bit too "broad," sometimes retrieving documents that were only loosely related to the topic.

6. Conclusion

Building a custom search engine allowed us to surpass the performance of a generic professional setup. For a scientific dataset like Cranfield, the best possible configuration is a Custom Index using Stemming and the TF-IDF or BM25 models. This setup provided a 10-15% accuracy improvement over basic search methods and proved capable of finding highly relevant research abstracts within the top 5 results.